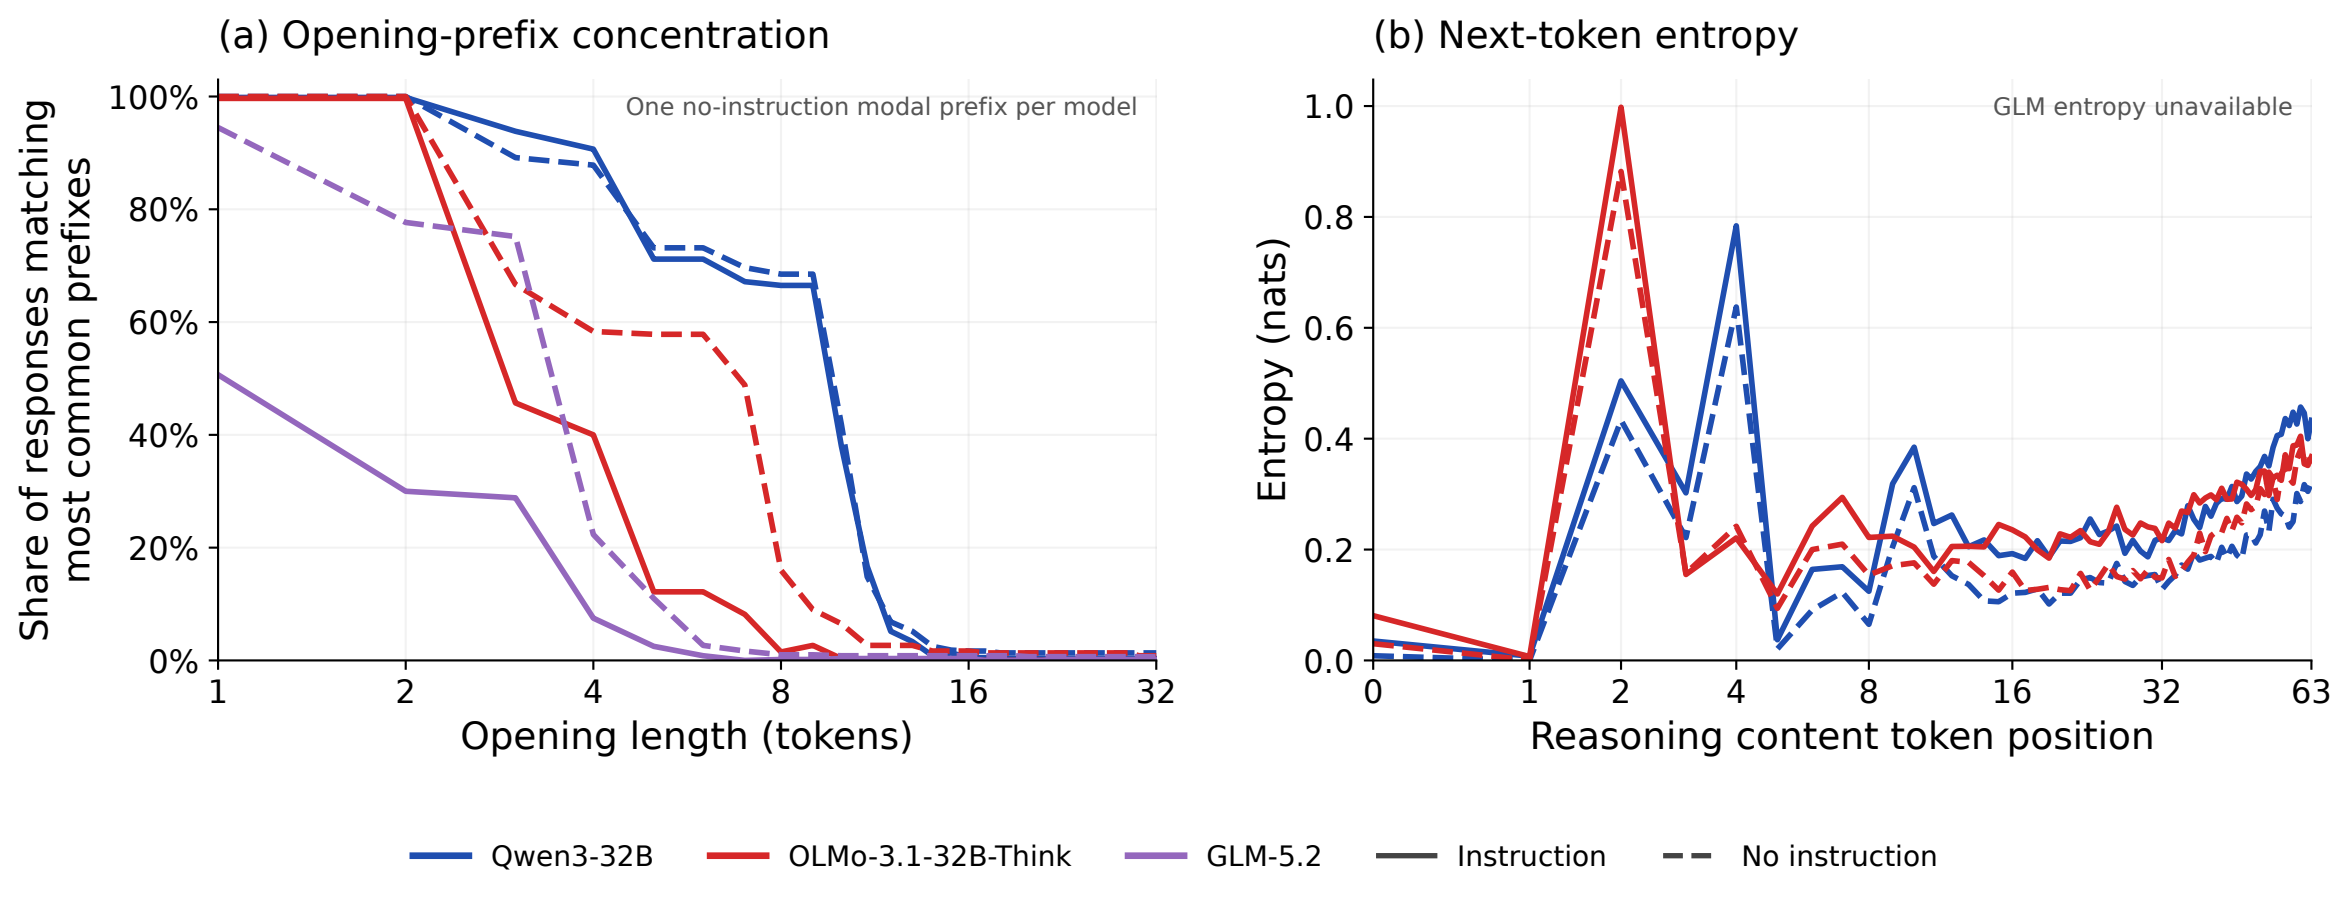
(a) Opening-prefix concentration
One no-instruction modal prefix per model
Share of responses matching most common prefixes
100%
80%
60%
40%
20%
0%
1
2
4
8
16
32
Opening length (tokens)
(b) Next-token entropy
GLM entropy unavailable
Entropy (nats)
1.0
0.8
0.6
0.4
0.2
0.0
0
1
2
4
8
16
32
63
Reasoning content token position
Qwen3-32B
OLMo-3.1-32B-Think
GLM-5.2
Instruction
No instruction